

Scheduler


Scheduler – in most simple terms

- **Selects a process to execute and passes control to it !**
 - The process is chosen out of “READY” state processes
 - Saving of context of “earlier” process and loading of context of “next” process needs to happen
- **Questions**
 - What are the different scenarios in which a scheduler called ?
 - What are the intricacies of “passing control”
 - What is “context” ?

Steps in scheduling scheduling

- Suppose you want to switch from P1 to P2 on a timer interrupt
- P1 was doing
`F() { i++; j++; }`
- P2 was doing
`G() { x--; y++; }`
- P1 will experience a timer interrupt, switch to kernel (scheduler) and scheduler will schedule P2

4 stacks need to change!



- **User stack of process -> kernel stack of process**
 - Switch to kernel stack
 - The normal sequence on any interrupt !
- **Kernel stack of process -> kernel stack of scheduler**
 - Why?
- **Kernel stack of scheduler -> kernel stack of new process . Why?**
- **Kernel stack of new process -> user stack of new process**

scheduler()

- **Disable interrupts**
- **Find a RUNNABLE process. Simple round-robin!**
- **c->proc = p**
- **switchvm(p) : Save TSS of scheduler's stack and make CR3 to point to new process pagedir**
- **p->state = RUNNING**
- **swtch(&(c->scheduler), p->context)**

swtch

swtch:

movl 4(%esp), %eax

movl 8(%esp), %edx

Save old callee-saved registers

pushl %ebp

pushl %ebx

pushl %esi

pushl %edi

Switch stacks

movl %esp, (%eax)

movl %edx, %esp

Load new callee-saved registers

popl %edi

popl %esi

popl %ebx

popl %ebp

ret

scheduler()

- **swtch(&(c->scheduler), p->context)**
- **Note that when scheduler() was called, when P1 was running**
- **After call to swtch() shown above**
 - **The call does NOT return!**
 - **The new process P2 given by 'p' starts running !**
 - **Let's review swtch() again**

switch(old, new)

- The magic function in switch.S
- Saves callee-save registers of old context
- Switches esp to new-context's stack
- Pop callee-save registers from new context

ret

- where? in the case of first process – returns to forkret() because stack was setup like that !
- in case of other processes, return where?
 - Return address given on kernel stack. But what's that?
 - The EIP in p->context
 - When was EIP set in p->context ?

scheduler()

- **Called from?**
 - `mpmain()`
 - No where else!
- **`sched()` is another scheduler function !**
 - Who calls `sched()` ?
 - `exit()` - a process exiting calls `sched ()`
 - `yield()` - a process interrupted by timer calls `yield()`
 - `sleep()` - a process going to wait calls `sleep()`

sched()

```
void
sched(void)
{
    int intena;
    struct proc *p = myproc();
    if(!holding(&ptable.lock))
        panic("sched ptable.lock");
    if(mycpu()->ncli != 1)
        panic("sched locks");
    if(p->state == RUNNING)
        panic("sched running");
    if(readeflags() & FL_IF)
        panic("sched interruptible");
    intena = mycpu()->intena;
    swtch(&p->context, mycpu()-
    >scheduler);
    /*A*/ mycpu()->intena = intena;
}
```

- **get current process**
- **Error checking code (ignore as of now)**
- **get interrupt enabled status on current CPU (ignore as of now)**
- **call to swtch**
 - **Note the arguments' order**
 - **p->context first, mycpu()->scheduler second**
- **swtch() is a function call**
 - **pushes address of /*A*/ on stack of current process p**
 - **switches stack to mycpu()->scheduler. Then pops EIP from that stack and jumps there.**
 - **when was mycpu()->scheduler set? Ans: during scheduler()!**

sched() and scheduler()

```
sched() {
```

```
...
```

```
    swtch(&p->context, mycpu()-  
>scheduler); /* X */
```

```
}
```

```
scheduler(void) {
```

```
...
```

```
    swtch(&(c->scheduler), p-  
>context); /* Y */
```

```
}
```

- scheduler() saves context in c->scheduler, sched() saves context in p->context
- after swtch() call in sched(), the control jumps to Y in scheduler
 - Switch from process stack to scheduler's stack
- after swtch() call in scheduler(), the control jumps to X in sched()
 - Switch from scheduler's stack to new process's stack
- Set of co-operating functions

sched() and scheduler() as co-routines

- **In sched()**
`swtch(&p->context, mycpu()->scheduler);`
- **In scheduler()**
`swtch(&(c->scheduler), p->context);`
- **These two keep switching between processes**
- **These two functions work together to achieve scheduling**
- **Using asynchronous jumps**
- **Hence they are co-routines**

To summarize

- **On a timer interrupt during P1**
 - trap() is called. **Stack has changed from P1's user stack to P1's kernel stack**
 - trap()->yield()
 - yield()->sched()
 - sched() -> swtch(&p->context, c->scheduler())
 - **Stack changes to scheduler's kernel stack.**
 - Switches to location "Y" in scheduler().
- **Now the loop in scheduler()**
 - calls switchkvm()
 - Then continues to find next process (P2) to run
 - Then calls switchvm(p): changing the page table to the P2's page tables
 - then calls swtch(&c->scheduler, p2's->context)
 - **Stack changes to P2's kernel stack.**
 - P2 runs the last instruction it was was in ! Where was it?
 - mycpu()->intena = intena; in sched()
 - Then returns to the one who called sched() i.e. exit/sleep, etc
 - Finally returns from it's own "TRAP" handler and returns to P2's user stack and user code