

Memory Management – Continued

More on Linking, Loading, Paging

Review of last class

- **MMU : Hardware features for MM**
- **OS: Sets up MMU for a process, then schedules process**
- **Compiler : Generates object code for a particular OS + MMU architecture**
- **MMU: Detects memory violations and raises interrupt --> Effectively passing control to OS**

More on Linking and Loading

- **Static Linking:** All object code combined at link time and a big object code file is created
- **Static Loading:** All the code is loaded in memory at the time of `exec()`
- **Problem:** Big executable files, need to load functions even if they do not execute
- **Solution:** Dynamic Linking and Dynamic Loading

Dynamic Linking

- Linker is normally invoked as a part of compilation process
 - Links
 - function code to function calls
 - references to global variables with “extern” declarations
- Dynamic Linker
 - Does not combine function code with the object code file
 - Instead introduces a “stub” code that is indirect reference to actual code
 - At the time of “loading” (or executing!) the program in memory, the “link-loader” (part of OS!) will pick up the relevant code from the library machine code file (e.g. libc.so.6)

Dynamic Linking on Linux

```
#include <stdio.h>

int main() {
    int a, b;
    scanf("%d%d", &a, &b);
    printf("%d %d\n", a, b);
    return 0;
}
```

PLT: Procedure Linkage Table

used to call external procedures/functions whose address is to be resolved by the dynamic linker at run time.

Output of `objdump -x -D`

Disassembly of section `.text`:

00000000000001189 <main>:

11d4: callq 1080 <printf@plt>

Disassembly of section `.plt.got`:

00000000000001080 <printf@plt>:

1080: endbr64

1084: bnd jmpq *0x2f3d(%rip) # 3fc8
<printf@GLIBC_2.2.5>

108b: nopl 0x0(%rax,%rax,1)

Dynamic Loading

- **Loader**
 - Loads the program in memory
 - Part of `exec()` code
 - Needs to understand the format of the executable file (e.g. the ELF format)
- **Dynamic Loading**
 - Load a part from the ELF file only if needed during execution
 - Delayed loading
 - Needs a more sophisticated memory management by operating system – to be seen during this series of lectures

Dynamic Linking, Loading

- **Dynamic linking necessarily demands an advanced type of loader that understands dynamic linking**
 - **Hence called ‘link-loader’**
 - **Static or dynamic loading is still a choice**
- **Question: which of the MMU options will allow for which type of linking, loading ?**

Continuous memory management

What is Continuous memory management?

- Entire process is hosted as one continuous chunk in RAM
- Memory is typically divided into two partitions
 - One for OS and other for processes
 - OS most typically located in “high memory” addresses, because interrupt vectors map to that location (Linux, Windows) !

Hardware support needed: base + limit (or relocation + limit)

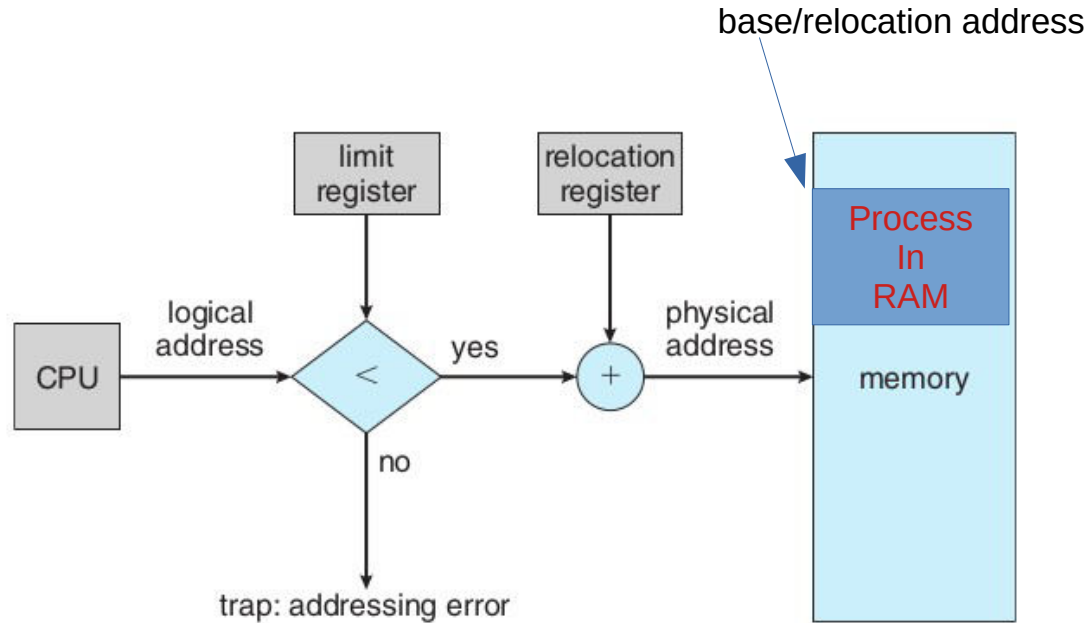


Figure 9.6 Hardware support for relocation and limit registers.

Problems faced by OS

- Find a continuous chunk for the process being forked
- Different processes are of different sizes
 - Allocate a size parameter in the PCB
- After a process is over – free the memory occupied by it
- Maintain a list of free areas, and occupied areas
 - Can be done using an array, or linked list

Variable partition scheme




Figure 9.7 Variable partition.

Problem: how to find a “hole” to fit in new process

- **Suppose there are 3 free memory regions of sizes 30k, 40k, 20k**
- **The newly created process (during fork() + exec()) needs 15k**
- **Which region to allocate to it ?**

Strategies for finding a free chunk


- 6k, 17k, 16k, 40k holes . Need 15k.
- **Best fit:** Find the smallest hole, larger than process. **Ans: 16k**
- **Worst fit:** Find the largest hole. **Ans: 40k**
- **First fit:** Find the “first” hole larger than the process. **Ans: 17k**

Problem : External fragmentation

- Free chunks: 30k, 40k, 20k
- The newly created process (during `fork()` + `exec()`) needs 50k
- Total free memory: $30+40+20 = 90k$
 - But can't allocate 50k !

Solution to external fragmentation

- **Compaction !**
- OS moves the process chunks in memory to make available continuous memory region
 - Then it must update the memory management information in PCB (e.g. base of the process) of each process
- Time consuming
- Possible only if the relocation+limit scheme of MMU is available



Another solution to external fragmentation: Fixed size partitions

- Fixed partition scheme
- Memory is divided by OS into chunks of equal size: e.g., say, 50k
 - If total 1M memory, then 20 such chunks
- Allocate one or more chunks to a process, such that the total size is \geq the size of the process
 - E.g. if request is 50k, allocate 1 chunk
 - If request is 40k, still allocate 1 chunk
 - If request is 60k, then allocate 2 chunks
- Leads to internal fragmentation
 - space wasted in the case of 40k or 60k requests above

50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K

Kernel
Kernel
Free
P1 (50 KB)
P2 (80 KB)
Unused (20 KB)
Free
P3 (120 KB)
Unused (30 KB)
Free
Free

Fixed partition scheme

- OS needs to keep track of
 - Which partition is free and which is used by which process
 - Free partitions can simply be tracked using a bitmap or a list of numbers
 - Each process's PCB will contain list of partitions allocated to it

Solution to internal fragmentation


- Reduce the size of the fixed sized partition
- How small then ?
 - Smaller partitions mean more overhead for the operating system in allocating deallocating

Paging

An extended version of fixed size partitions

- **Partition = page**
 - **Process** = logically continuous sequence of bytes, divided in 'page' sizes
 - **Memory** divided into equally sized page 'frames'
- **Important distinction**
 - **Process** need not be continuous in RAM
 - **Different** page sized chunks of process can go in any page frame
 - **Page** table to map pages into frames

Logical address seen as



Paging hardware




Figure 9.8 Paging hardware.

MMU's job

To translate a logical address generated by the CPU to a physical address:

1. Extract the page number p and use it as an index into the page table.

(Page table location is stored in a hardware register

Also stored in PCB of the process, so that it can be used to load the hardware register on a context switch)

2. Extract the corresponding frame number f from the page table.

3. Replace the page number p in the logical address with the frame number f .

Job of OS

- **Allocate a page table for the process, at time of fork()/exec()**
 - **Allocate frames to process**
 - **Fill in page table entries**
- **In PCB of each process, maintain**
 - **Page table location (address)**
 - **List of pages frames allocated to this process**
- **During context switch of the process, load the PTBR using the PCB**

Job of OS

- **Maintain a list of all page frames**
 - Allocated frames
 - Free Frames (called frame table)
 - Can be done using simple linked list
 - Innovative data structures can also be used to maintain free and allocated frames list (e.g. xv6 code)
 -





Figure 9.11 Free frames (a) before allocation and (b) after allocation.

Disadvantage of Paging

- **Each memory access results in two memory accesses!**
 - One for page table, and one for the actual memory location !
 - Done as part of execution of instruction in hardware (not by OS!)
 - Slow down by 50%

Speeding up paging

- Translation Lookaside Buffer (TLB)
- Part of CPU hardware
- A cache of Page table entries
- Searched in parallel for a page number



Speedup due to TLB

- Hit ratio
- Effective memory access time
 - = Hit ratio * 1 memory access time + miss ratio * 2 memory access time
- Example: memory access time 10ns, hit ratio = 0.8, then
 - effective access time = $0.80 \times 10 + 0.20 \times 20$
 - = 12 nanoseconds

12,287

10,468

page 5
page 4
page 3
page 2
page 1
page 0

00000

frame number

valid-invalid bit

0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i


page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page n

**Memory
protection
with paging**

Figure 9.13 Valid (v) or invalid (i) bit in a page table.

X86 PDE and PTE




PDE

- P Present
- W Writable
- U User
- WT 1=Write-through, 0=Write-back
- CD Cache disabled
- A Accessed
- D Dirty
- PS Page size (0=4KB, 1=4MB)
- PAT Page table attribute index
- G Global page
- AVL Available for system use



PTE



**Shared
pages (e.g.
library)
with paging**

Figure 9.14 Sharing of standard C library in a paging environment.

Paging: problem of large PT

- 64 bit address
- Suppose 20 bit offset
 - That means $2^{20} = 1 \text{ MB}$ pages
 - 44 bit page number: 2^{44} that is trillion sized page table!
 - Can't have that big continuous page table!

Paging: problem of large PT

- 32 bit address
- Suppose 12 bit offset
 - That means $2^{12} = 4 \text{ KB}$ pages
 - 20 bit page number: 2^{20} that is a million entries
 - Can't always have that big continuous page table as well, for each process!

Hierarchical paging





Figure 9.15 A two-level page-table scheme.

outer page	inner page	offset
p_1	p_2	d
42	10	12

More hierarchy

2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

Problems with hierarchical paging

- More number of memory accesses with each level !
 - Too slow !
- OS data structures also needed in that proportion

Hashed page table




Figure 9.17 Hashed page table.

Inverted page table

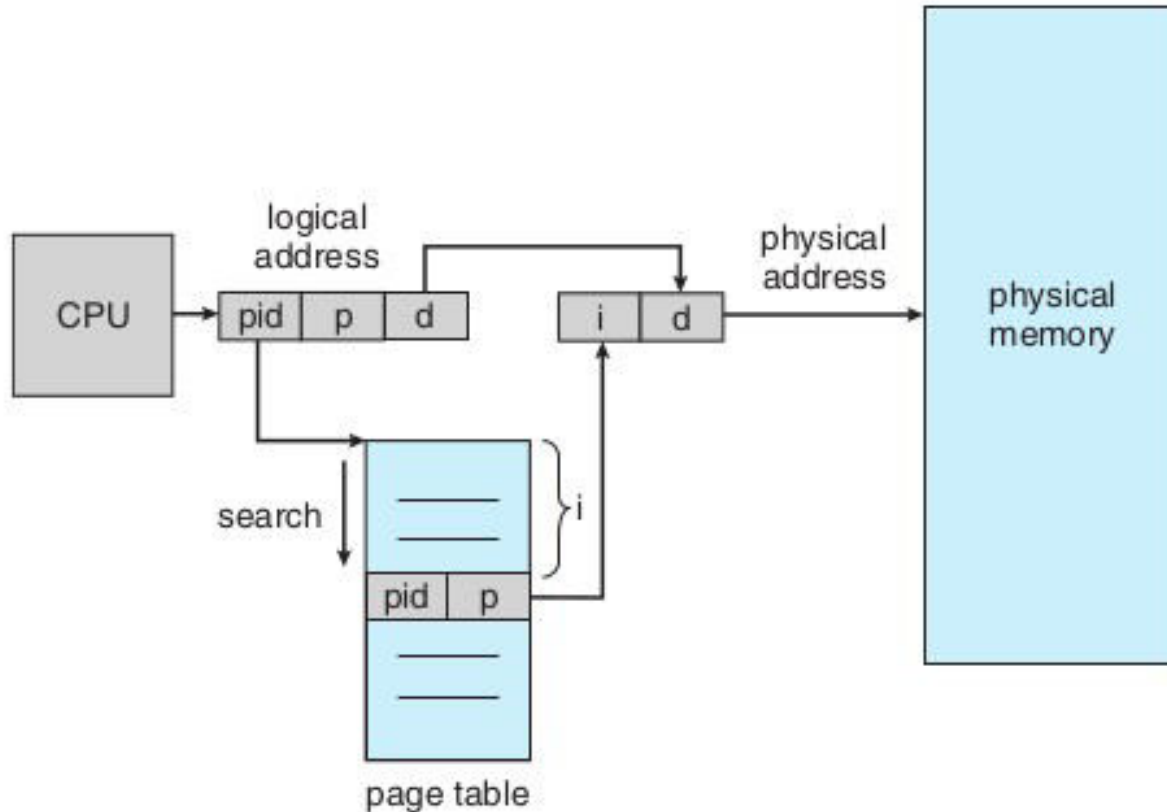


Figure 9.18 Inverted page table.

Normal page table – one per process --> Too much memory consumed

Inverted page table : global table – only one
Needs to store PID in the table entry

Examples of systems using inverted page tables include the 64-bit Ultra SPARC and Power PC

virtual address
consists of a triple:
<process-id, page-number, offset>

Case Study: Oracle SPARC Solaris

- 64 bit SPARC processor , 64 bit Solaris OS
- Uses Hashed page tables
 - one for the kernel and one for all user processes.
 - Each hash-table entry : $\text{base} + \text{span} (\# \text{pages})$
 - Reduces number of entries required

Case Study: Oracle SPARC Solaris

- **Caching levels: TLB (on CPU), TSB(in Memory), Page Tables (in Memory)**
 - CPU implements a TLB that holds translation table entries (TTE s) for fast hardware lookups.
 - A cache of these TTEs resides in a in-memory translation storage buffer (TSB), which includes an entry per recently accessed page
 - When a virtual address reference occurs, the hardware searches the TLB for a translation.
 - If none is found, the hardware walks through the in memory TSB looking for the TTE that corresponds to the virtual address that caused the lookup

Case Study: Oracle SPARC Solaris

- If a match is found in the TSB , the CPU copies the TSB entry into the TLB , and the memory translation completes.
- If no match is found in the TSB , the kernel is interrupted to search the hash table.
- The kernel then creates a TTE from the appropriate hash table and stores it in the TSB for automatic loading into the TLB by the CPU memory-management unit.
- Finally, the interrupt handler returns control to the MMU , which completes the address translation and retrieves the requested byte or word from main memory.

Swapping




Figure 9.19 Standard swapping of two processes using a disk as a backing store.

Swapping

- **Standard swapping**
 - Entire process swapped in or swapped out
 - With continuous memory management
- **Swapping with paging**
 - Some pages are “paged out” and some “paged in”
 - Term “paging” refers to paging with swapping now




Figure 9.20 Swapping with paging.

Words of caution about 'paging'

- Not as simple as it sounds when it comes to implementation
 - Writing OS code for this is challenging